

Why Kubernetes Exists

Kubernetes is an open-source platform for hosting containers and providing a cloud-neutral API to manage storage, networking, security, and other resources for applications.

- Provides **continuous reconciliation** of the desired state of deployments.
- Enables **reproducibility** of infrastructure, eliminating drift across identical components.
- Integrates **DevOps** practices directly into the application lifecycle.

Key Terms Review

A quick reference for essential Kubernetes terminology.

- **CNI / CSI** – Container Networking Interface / Container Storage Interface, enabling pluggable networking and storage.
- **Container** – Docker or OCI image that runs an application.
- **Control plane** – “Brains” of the cluster (API server, scheduler, controller manager).
- **DaemonSet** – Runs a copy of a Pod on **every** node.
- **Deployment** – Manages a set of identical Pods.
- **kubectl** – CLI for interacting with the Kubernetes API server.
- **kubelet** – Node-side agent that enforces the desired state.
- **Node** – A machine (VM or physical) that runs a kubelet.
- **OCI** – Open Container Initiative image format.
- **Pod** – Smallest deployable unit; encapsulates one or more containers.

Infrastructure Drift Problem

Managing infrastructure as code prevents “drift” where servers diverge from their intended configuration.

Common manual toil avoided by Kubernetes:

| Manual Task | Why It Causes Drift | Kubernetes Solution |
|---------------------------------------|--------------------------------------|-----------------------------------|
| Updating Java version on many servers | Missed hosts → inconsistent runtimes | Declarative Pod/Deployment specs |
| Scaling hardware manually | Human error, delayed provisioning | Autoscaling controllers |
| Hand-crafting load-balancer rules | Inconsistent routing | Service objects + kube-proxy |
| Forgetting to document changes | Knowledge loss | Version-controlled YAML manifests |

Containers & Images

Containers add isolation, allowing multiple versions of the same library to coexist.

```
# Dockerfile for a MySQL image
FROM alpine:3.15.4
RUN apk add --no-cache mysql
ENTRYPOINT ["/usr/bin/mysqld"]
```

- Images are **OCI tarballs** with layered filesystem contents.
- Running a container unpacks the image and starts a process on the host.
- Example of running two containers manually:

```
docker run -t -i ui -p 80:80
docker run -t -i microservice-a -p 8080:8080
```

```
docker run -t -i microservice-b -p 8081:8081  
docker run -t -i cockroach:cockroach -p 6379:6379
```


🛠 Core Foundation of Kubernetes

- All objects are defined as [plain YAML](#) (or JSON) files.
- The same syntax configures Pods, Services, RBAC, Storage, etc.
- Example: building a simple MySQL pod.

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: core-mysql  
spec:  
  containers:  
    - name: mysql  
      image: myregistry.com/mysql-server:v1.0
```

```
kubectl create -f core-mysql.yaml
```

💻 Kubernetes Architecture & Components



The diagram visualizes the control plane (API server, scheduler, controller manager) communicating with multiple worker nodes, each running a kubelet.

Key components:

- **Control plane** – API server, scheduler, controller manager, etcd.
- **Worker nodes** – Run kubelet, container runtime, kube-proxy, and a CNI plugin.
- **etcd** – Consistent key-value store for cluster state.

The Kubernetes API

The API server exposes a **RESTful** interface for CRUD operations on all cluster objects.

- Objects have **apiVersion**, **kind**, and **metadata**.
- Example of listing API resources:

```
kubectl api-resources | head
```

Output snippet:

```
NAME      SHORTNAMES   NAMESPACED   KIND
bindings    ...          false        Binding
componentstatuses  cs       false        ComponentStatus
configmaps    cm          true         ConfigMap
...
```

Example Use Cases

1 Online Retailer

- Needs rapid seasonal scaling.
- Benefits from Kubernetes auto-scaling, high availability, and cloud-agnostic deployment.

2 Online Giving Platform

- Grew from a single VM to many microservices (Python, Java, NGINX, etc.).
- Migration to Kubernetes reduced VM count from ~30 to 5, eliminated Puppet, and improved disaster recovery.

? When Not to Use Kubernetes

| Scenario | Reason |
|--|---|
| High-Performance Computing (HPC) | Container overhead can add latency; nanosecond precision may be required. |
| Legacy monoliths with strict hardware dependencies | Hard to containerize or unsupported by vendors. |
| Rigid legacy migrations | Minimal benefit if the app cannot be broken into microservices. |

The Pod Concept

A **Pod** is the smallest Kubernetes object that encapsulates one or more containers sharing Linux namespaces.


- Provides a **network namespace**, **PID namespace**, **IPC namespace**, **cgroup**, **mount**, and **user** namespaces.
- Pods are created via higher-level controllers (Deployments, StatefulSets, DaemonSets).

Pod YAML Example

```
apiVersion: v1
kind: Pod
metadata:
  name: busybox-pod
spec:
  containers:
    - name: busybox
      image: busybox
      command: ["sleep", "3600"]
```

```
kubectl apply -f busybox-pod.yaml
```

🌐 Example Web Application (Zeus Zap)



The diagram breaks the system into UI, services, and infrastructure layers, illustrating how microservices communicate.

Components:

- **UI** – NGINX serving a single-page app.

- **Microservices** – Two Python/Django services.
- **Database** – CockroachDB (PostgreSQL-compatible) on port 6379.
- **Messaging, caching, search** – Additional services (Redis, Elasticsearch, etc.).

Operational Challenges Highlighted

- **Port conflicts** when running multiple UI containers.
- **Dynamic IP handling** for database relocation.
- **High availability** for the database.
- **Burst traffic** from e-sport events.

Infrastructure for Web Application

To run the above stack without orchestration, you'd need:

- Multiple VMs or physical servers.
- Load balancers, DNS, persistent storage, security policies, CI/CD pipelines, monitoring, and logging.

Kubernetes abstracts all these via declarative objects, reducing manual toil.

Pod Definition & Creation

```
apiVersion: v1
kind: Pod
metadata:
  name: zeus-ui
spec:
  containers:
    - name: nginx
      image: nginx:1.21
      ports:
        - containerPort: 80
```

```
kubectl apply -f zeus-ui.yaml
```

- The **kubelet** on the chosen node watches the API for this Pod, pulls the image via the CRI, and starts it.
- A **pause container** is launched first to set up the network namespace.

Linux Namespaces in Pods

| Namespace | Purpose |
|----------------|---|
| PID | Isolates process IDs. |
| Network | Gives each Pod its own IP address. |
| IPC | Isolates inter-process communication (shared memory, semaphores). |
| cgroup | Controls CPU, memory, and I/O limits. |
| Mount | Provides a separate filesystem view. |
| User | Maps user IDs inside the container. |

These namespaces enable the Pod to behave like a lightweight VM while sharing the host kernel.

Node API Object

A **Node** is a Kubernetes API object representing a worker machine.

```

apiVersion: v1
kind: Node
metadata:
  name: kind-control-plane
  labels:
    kubernetes.io/hostname: kind-control-plane
spec:
  podCIDR: 10.244.0.0/24

```

```

apiVersion: v1
kind: Node
metadata:
  annotations:
    kubeadm.alpha.kubernetes.io/cri-socket:
      /run/containerd/containerd.sock

```

The CRI socket used. With kind (and most clusters), this is the containerd socket.

The snippet shows a Node's metadata, spec, and status sections, including its CIDR, capacity, and condition fields.

Key fields:

- **metadata.annotations** – e.g., node.alpha.kubernetes.io/ttl.
- **spec.podCIDR** – IP range for Pods on this node.
- **status.conditions** – MemoryPressure, DiskPressure, PIDPressure, Ready, etc.

Control Plane & Scaling

- The **scheduler** assigns Pods to Nodes based on resource availability, affinity/anti-affinity, and taints.
- **Controller manager** runs loops that ensure desired state (e.g., replica count, StatefulSet ordering).
- **Autoscaling** options:
 1. **Horizontal Pod Autoscaler (HPA)** – Adds/removes Pods.
 2. **Vertical Pod Autoscaler (VPA)** – Adjusts CPU/memory requests.
 3. **Cluster Autoscaler** – Adds/removes Nodes.

Scaling Example

```
# Increase UI deployment to 300 replicas
kubectl scale --replicas=300 deployment zeus-front-end-ui
```

Sequence of events (illustrated in the book):

1. kubectl scale updates the Deployment's spec.
2. The **ReplicationController** notices the change.
3. The **scheduler** picks Nodes for new Pods.
4. The **kubelet** on each Node creates the Pods via the CRI.

Autoscaling & Cost Management

- **Pod density** – Packing many Pods onto a node reduces infrastructure cost.
- Steps to achieve optimal density:
 1. **Profile** each application's CPU/memory usage.
 2. **Select** node sizes that can host multiple Pods comfortably.
 3. **Group** compatible workloads using **taints**, **tolerations**, and **affinity** rules.
- **Noisy neighbor** mitigation: use **PodPriority**, **resource quotas**, and **QoS classes**.

Cost-saving tip: Shut down development clusters during weekends (set node count to zero) and spin them up on demand.

Summary Highlights

- **Pod** = atomic deployment unit built from Linux namespaces.

- **Control plane** orchestrates Pods, storage, networking, and scaling.
- **kubectl** is the single tool for all CRUD operations against the API server.
- **Declarative YAML** drives the entire cluster, from Nodes to Services.
- **Autoscaling** and **pod density** are core mechanisms for balancing performance with cost.

---##  Fundamental Linux Primitives in Kubernetes

- **kube-proxy** – creates *iptables* rules that route Service traffic to Pods.
- **Container Network Interface (CNI)** – uses the same network proxy for implementing **NetworkPolicies**.
- **Container Storage Interface (CSI)** – socket used by the *kubelet* to talk to storage back-ends (e.g., EBS, NFS).
- **Container runtime commands** – unshare, mount, nsenter, kill, ps, ls etc. are the building blocks for isolation, networking and process management.

Definition – *Linux primitives* are low-level OS tools that manipulate files, devices, sockets or namespaces; Kubernetes composes them to implement Pods, Services, and other resources.

Key Commands Table

| Command | Primary Use | Security Implication |
|----------|---|---|
| swapoff | Disable swap (required for deterministic CPU/Memory accounting) | Prevents memory paging that could hide OOM events |
| iptables | Set up network forwarding / firewall rules (used by kube-proxy) | Misconfiguration can expose services |
| mount | Bind a device or directory into a namespace | Improper binds may leak host data |
| systemd | Starts the <i>kubelet</i> process | Central to node lifecycle |
| socat | Bidirectional stream forwarding (used by kubectl port-forward) | Allows remote access to pods |
| nsenter | Enter another process's namespaces (PID, network, mount) | Enables deep debugging |
| unshare | Create new namespaces for a child process (PID, net, mount) | Provides strong isolation |
| ps | List processes (used to detect zombies, rogue pods) | Visibility into container health |

Building and Running Pods

Minimal BusyBox Pod YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: core-k8s
  labels:
    role: just-an-example
    app: my-example-app
spec:
  containers:
  - name: busybox
    image: docker.io/busybox:latest
    command: ['sleep','10000']
    ports:
    - name: webapp-port
      containerPort: 80
      protocol: TCP
```

- **Labels** enable selection by Services or kubectl queries.
- The **command** runs a long-running sleep so the Pod stays alive for inspection.

Inspecting the Pod

```
# List all Pods (all namespaces)
kubectl get pods --all-namespaces

# Show only the Pod phase
kubectl get pods -o=jsonpath='{.items[0].status.phase}'
# → Running

# Show Pod IP
kubectl get pods -o=jsonpath='{.items[0].status.podIP}'
# → 10.244.0.11

# Show host IP
kubectl get pods -o=jsonpath='{.items[0].status.hostIP}'
# → 172.17.0.2
```

Linux Commands & Utilities for Kubernetes

| Utility | Role in Pod Lifecycle |
|----------|--|
| swapoff | Guarantees that memory limits are respected. |
| iptables | Implements Service routing (kube-proxy). |
| mount | Exposes storage volumes inside the container. |
| systemd | Starts the kubelet, which drives the control loop. |
| socat | Powers kubectl port-forward. |
| nsenter | Enters a Pod's namespaces for debugging. |
| unshare | Creates isolated PID, network, and mount namespaces. |
| ps | Checks for zombie processes or runaway containers. |

“Everything Is a File” (or File Descriptor)

Quote – “*Everything is a file*” is a core Linux principle; devices, sockets, and even process information appear as files under /dev, /proc, or /sys.

- **Directories** are files containing lists of other file names.
- **Device files** (/dev/eth0) let tools like ls verify device presence.
- **Sockets & pipes** appear as special files enabling inter-process communication.

The **/proc** filesystem provides real-time process metadata, e.g., /proc/<PID>/cgroup.

Using chroot to Isolate a Process

The following script creates a minimal sandbox (a “poor-man’s VM”) and drops into a Bash shell:

```
#!/bin/bash
mkdir -p /home/namespace/box/{bin,lib,lib64,proc}
cp -v /usr/bin/kill /home/namespace/box/bin/
cp -v /usr/bin/ps   /home/namespace/box/bin/
```


```

cp -v /bin/bash    /home/namespace/box/bin/
cp -v /bin/ls      /home/namespace/box/bin/
cp -r /lib/*       /home/namespace/box/lib/
cp -r /lib64/*     /home/namespace/box/lib64/
mount -t proc proc /home/namespace/box/proc
chroot /home/namespace/box /bin/bash

```

- **Result:** The Bash process sees only the files explicitly copied into /home/namespace/box.
- **Use case:** Safe testing of destructive commands (rm -rf /) without harming the host.

Figure 3.3 – chroot vs Host Filesystem



The diagram shows how the chrooted environment mounts only a subset of the host's directories, illustrating the isolation achieved by chroot.

Using mount for Data Access

```

# Bind-mount /tmp into the sandbox as /data
mount --bind /tmp /home/namespace/box/data

```

- The sandbox can now read/write /data/*, which maps to the host's /tmp.
- **Security note:** Bind-mounts expose host data; Kubernetes often disables hostPath in production for this reason.

🔒 Securing Processes with unshare

Running the previous chroot script inside an unshare namespace prevents the sandbox from seeing or killing host processes (e.g., the kubelet).

```

unshare -p -n -f --mount-proc=/home/namespace/box/proc \
        chroot /home/namespace/box /bin/bash

```

- **-p** – new PID namespace (the process thinks it is PID 1).
- **-n** – new network namespace (isolated lo interface).
- **-f** – fork after unsharing.

Figure 3.4 – unshare Process Isolation

```
# Note that this won't work if you've already unmounted this directory.  
root@kcp:/# unshare -p -f  
    --mount-proc=/home/namespace/box/proc  
    chroot /home/namespace/box /bin/bash      ↪ Creates a new shell running  
                                                in a namespace  
  
bash-5.0# ps -ax  
PID TTY      STAT      TIME COMMAND      ↪ Observes all processes visible to the  
1 ?          S          0:00 /bin/bash  
2 ?          R+         0:00 ps -ax           namespaces; seems kind of low, right?
```

The output demonstrates a tiny process table visible inside the new namespace, confirming isolation from host processes.

Creating a Network Namespace

```
unshare -p -n -f --mount-proc=/home/namespace/box/proc \  
    chroot /home/namespace/box /bin/bash  
# Inside the shell  
ip a    # Shows only lo and any manually added interfaces
```

- No eth0 is present because a CNI plugin has not been invoked.
- Real Pods receive an eth0 with an IP from the cluster's overlay network (e.g., 10.244.0.x).

Inspecting Pods and Services

Using iptables to View Service Rules

```
# Show all iptables rules that reference cluster IPs  
iptables-save | grep -E 'KUBE-SVC|KUBE-SEP'
```

- **KUBE-SVC...** rules forward traffic from a Service IP to one of its endpoints.
- **KUBE-SEP...** rules perform DNAT to the Pod's IP and port.

kube-dns Example

```
kubectl exec -ti core-k8s -- mount | grep resolv.conf  
# → /dev/sdal on /etc/resolv.conf type ext4 (rw,relatime)
```

- The DNS configuration is mounted from the host, enabling Pods to resolve service names.

Kind: Local Kubernetes Cluster Setup

1. Install **Docker**, **kubectl**, and **kind**.
2. Create a cluster:


```
kind delete cluster --name=kind    # Clean up any previous cluster  
kind create cluster
```

3. Verify:

```
kubectl get pods --all-namespaces  
docker ps    # Shows the kind node container
```

```
docker exec -it <container-id> /bin/sh # SSH-like access to the node
```

Figure 4.1 – Kind Cluster Workflow



Kind builds a node inside a Docker container, enabling rapid local experimentation.

✓ cgroups: Resource Management

- **cgroups** (control groups) partition CPU, memory, I/O, etc. among processes.
- Kubernetes creates a cgroup per Pod under `/sys/fs/cgroup/.../kubepods/...`.

Example: Limiting Memory to 10 bytes

```
mkdir /sys/fs/cgroup/memory/chroot0
echo "10" > /sys/fs/cgroup/memory/chroot0/memory.limit_in_bytes
```

```
echo "0" > /sys/fs/cgroup/memory/chroot0/memory.swappiness
echo $$ > /sys/fs/cgroup/memory/chroot0/tasks # Add current PID
```

Running ls now fails with “*Cannot allocate memory*”.

Resources Stanza (Greedy Pod)

```
apiVersion: v1
kind: Pod
metadata:
  name: greedy-pod
spec:
  containers:
  - name: busybox
    image: busybox:latest
    command: ['sleep','10000']
    resources:
      limits:
        memory: "200Mi"
      requests:
        memory: "100Mi"
```

```
$ cat << EOF > greedy-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: core-k8s-greedy
spec:
  containers:
  - name: any-old-name-will-do
    image: docker.io/busybox:latest
    command: ['sleep','10000']
    resources:           ←
      limits:
        memory: "200Mi"   | Tells Kubernetes to create
      requests:          | a cgroup to limit (or not)
        memory: "100Mi"   | available CPU
EOF
```

The manifest tells the kubelet to create a cgroup that caps the container’s memory.

Allocatable Resources Equation

Formula – Allocatable = Node capacity – kube-reserved – system-reserved

- **Node capacity** – total CPU, memory, storage on the node.
- **kube-reserved** – resources needed by the kubelet, container runtime, etc.
- **system-reserved** – OS daemon overhead.

⌚ QoS Classes

| QoS Class | Definition | Typical Use |
|------------|---|---|
| Guaranteed | Both requests and limits are set for CPU and memory . | Critical services that must not be evicted. |

| | | |
|-------------------|--|---|
| Burstable | requests are set, but limits may be higher or omitted. | Workloads that need a baseline but can burst. |
| BestEffort | No requests or limits defined. | Low-priority jobs; first to be killed under pressure. |

Kubernetes automatically assigns the class based on the Pod's resources stanza.


Monitoring with Prometheus & cAdvisor

- **cAdvisor** (built into the kubelet) exposes per-cgroup metrics (CPU, memory, network).
- **Prometheus** scrapes these metrics from the API server's /metrics/cadvisor endpoint.

Prometheus Configuration (scrape every 3s)

```
global:
  scrape_interval: 3s
  evaluation_interval: 3s
scrape_configs:
  - job_name: prometheus
    metrics_path: /api/v1/nodes/kind-control-plane/proxy/metrics/cadvisor
    static_configs:
      - targets: ['localhost:8001']
```

Figure 4.2 – Prometheus Monitoring Architecture



The flowchart shows Prometheus pulling metrics from the API server, which aggregates data from kubelets (cAdvisor).

- **Metrics Types**
 - **Gauge** – current value (e.g., CPU usage).
 - **Counter** – monotonically increasing (e.g., total requests).
 - **Histogram** – distribution of latency or size.
- Example query to see etcd latency:

```
kubectl proxy &
curl localhost:8001/metrics | grep etcd_request_duration_seconds_bucket
```

📦 Summary of Core Concepts

- **Linux primitives** (iptables, mount, unshare, cgroups) are the scaffolding behind every Kubernetes feature.
- **Pods** are essentially isolated processes created with chroot, mount, and namespace isolation.

- **cgroups** enforce resource limits; QoS classes derive from the resources stanza.
- **Kind** provides a lightweight Docker-based test cluster; useful for hands-on experimentation.
- **Prometheus + cAdvisor** give deep visibility into kernel-level resource usage, enabling proactive troubleshooting. ## 

Key metrics to monitor in a busy cluster


| Metric | Description |
|------------------------------|--|
| container_memory_usage_bytes | Current memory consumption of a container. |
| container_fs_writes_total | Total number of filesystem write operations. |
| container_memory_cache | Amount of memory used for caching. |

Definition – *Gauge*: a metric that represents a single numerical value that can go up or down (e.g., cluster “up” flag).

Definition – *Histogram*: records the distribution of a set of values (e.g., request latency).

Definition – *Counter*: a cumulative metric that only increases (e.g., total requests served).

Figure 4.3 – Plotting metrics in a busy cluster



The line graph visualizes spikes in `container_cpu_user_seconds_total` after running a stress-test YAML. Hovering over the metrics links the values back to the specific containers and processes generating them.

Example: Generating Load

```
# WARNING: this creates heavy CPU and network load
kubectl apply -f https://raw.githubusercontent.com/giantswarm/kube-stresscheck/master/examples/node.yaml
```

Running the above on a laptop will cause noticeable CPU spikes and fan noise, illustrating how Prometheus captures rapid changes in container activity.

Kubernetes Service Types & SDN Basics

| Service Type | IP Scope | Typical Use |
|---------------------|----------------------------------|--|
| ClusterIP | Internal only | Service-to-service communication inside the cluster. |
| NodePort | Exposes on every node's IP | Simple external access; maps a random high port (e.g., 50491) to the service port. |
| LoadBalancer | External cloud IP (if supported) | Public-facing services; cloud provider provisions a load-balancer IP. |

Definition – *kube-proxy*: the component that programs kernel-level routing (iptables/IPVS) to forward Service IPs to the appropriate Pod IPs.

Service creation example (ClusterIP)

```
kubectl create service clusterip my-service-1 --tcp="100:100"
```


Repeated three times yields three distinct Service IPs in the 10.96.0.0/12 range, e.g.:

```
svc-1 ClusterIP 10.96.7.53 80/TCP
svc-2 ClusterIP 10.96.152.223 80/TCP
svc-3 ClusterIP 10.96.43.92 80/TCP
```

NodePort conversion – modify the Service YAML to set type: NodePort; Kubernetes allocates a random port (e.g., 30357) that forwards to the original Service port.

kube-proxy Routing Flow

Figure 5.1 – Traffic flow from a LoadBalancer into a Kubernetes cluster



The diagram shows how external traffic hits the LoadBalancer IP, is translated to a Service IP (ClusterIP), then routed by kube-proxy (iptables or IPVS) to the target Pod.

Key points

1. External IP (LoadBalancer) → Service IP (ClusterIP)

2. [kube-proxy](#) installs NAT rules (DNAT/SNAT) in the kernel.

3. Packets reach the Pod's IP on the node where the Pod lives.

kube-proxy operates in three modes: iptables, ipvs, and userspace. Most clusters default to iptables.

🛠️ CNI Specification & Provider Landscape

Definition – *CNI (Container Network Interface)*: a generic specification that defines how a container runtime requests network resources (add, delete, check).

CNI provider categories


| Provider | Primary Technology | Notable Features |
|----------------------------------|------------------------|--|
| Calico | BGP (Layer 3) | Supports IPIP, VXLAN, XDP; can replace kube-proxy. |
| Antrea | Open vSwitch (Layer 2) | Bridges pods via OVS; includes Antrea Proxy. |
| Flannel | Simple bridge | Legacy, minimal feature set. |
| Cilium | eBPF/XDP | High-performance, kernel-bypass networking. |
| KindNet | Minimal CNI for kind | Single-subnet, for local testing only. |
| Cloud-specific (GKE, EKS, Azure) | Proprietary | Tight integration with cloud load-balancers. |

The CNI [binary](#) implements three core commands:

1. [ADD](#) – attach a container to a network.
2. [DELETE](#) – detach a container.
3. [CHECK](#) – verify proper setup.

🔌 Calico vs. Antrea Architecture

Figure 5.2 – CNI networking in Calico and Antrea



Both plugins expose `add()`, `del()`, `check()` to the container runtime, then use different data-plane technologies (BGP vs. OVS) to program routes.

| Aspect | Calico | Antrea |
|------------------------------|-----------------------------|-----------------------------|
| Routing | Layer 3 BGP routes per node | Layer 2 OVS bridge per node |
| Tunnel types | IPIP, VXLAN, Geneve | Geneve (via OVS) |

| | | |
|--------------------------|----------------------------------|----------------------------|
| NetworkPolicy | iptables rules | OVS flow rules (table 90) |
| Proxy replacement | Can replace kube-proxy (via BGP) | Antrea Proxy (OVS-based) |
| Deployment | DaemonSet + controller pod | DaemonSet + controller pod |

Both use a **DaemonSet** to run a per-node agent and a **controller** pod to aggregate policies.

Installing Calico in a kind Cluster

1. [Disable the default kind-net CNI](#) (allows a real CNI).
2. [Create a custom kind config](#) (sets a large pod subnet).

```
# kind-Calico-conf.yaml
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
networking:
  disableDefaultCNI: true
  podSubnet: 192.168.0.0/16
nodes:
  - role: control-plane
  - role: worker
```

```
kind create cluster --name=calico --config=kind-Calico-conf.yaml
```

3. [Install Calico manifests](#)

```
wget https://docs.projectcalico.org/manifests/calico.yaml
kubectl apply -f calico.yaml
```

After installation, `kubectl get pods -n kube-system` should show `calico-node` and `calico-kube-controllers` pods in **Ready** state.

Verifying Calico routes

```
# On a Calico node
ip route
```

Typical output includes routes via tunl0 for remote node subnets and cali* interfaces for local pods.

Installing Antrea in a kind Cluster

```
# kind-Antrea-conf.yaml
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
networking:
  disableDefaultCNI: true
  podSubnet: 192.168.0.0/16
nodes:
  - role: control-plane
  - role: worker
```

```
kind create cluster --name=antrea --config=kind-Antrea-conf.yaml
kubectl apply -f https://github.com/vmware-tanzu/antrea/releases/download/v0.8.0/antrea.yaml -n kube-system
```

Antrea-specific interfaces

- antrea-gw0 – OVS gateway handling all pod traffic.
- ovs-system – OVS bridge managing the virtual switch.
- genev_sys_* – Geneve tunnel endpoints.

```
# Show Antrea OVS flows (policy table 90)
kubectl -n kube-system exec -it $(kubectl get pod -n kube-system -l app=antrea-agent -o name) -- ovs-ofctl du
```

CNI Routing & Tunnels

| CNI | Tunnel Device | Routing Style |
|---------|---------------------|--|
| Calico | tunl0 (IPIP) | Per-Pod route entries (many). |
| Antrea | antrea-gw0 (Geneve) | One gateway route per node (.1 address). |
| Flannel | flannel.1 (VXLAN) | Simple overlay, one route per node. |

In Calico, each remote pod appears as a separate route entry; Antrea collapses traffic to a single gateway address, simplifying the kernel routing table.

NetworkPolicy Implementation

Definition – *NetworkPolicy*: a Kubernetes API object that defines allow-list rules for pod traffic (ingress, egress, or both).

Simple deny-all policy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: web-deny-all
spec:
  podSelector:
    matchLabels:
      app: web
  ingress: [] # no allowed inbound traffic
```

Allow-only port 80 from app: web2

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: web-allow-http
spec:
  podSelector:
    matchLabels:
      app: web
  ingress:
  - ports:
    - port: 80
      from:
      - podSelector:
          matchLabels:
            app: web2
```

How policies appear in the data plane

- [Calico](#) – adds iptables chains (cali-tw-...) with MARK and DROP rules.
- [Antrea](#) – creates OVS flows in table 90, using conjunction rules to combine multiple criteria.

Inspecting Calico iptables changes

```
iptables-save | grep -A2 "cali-tw"
```

Inspecting Antrea OVS flows

```
kubectl -n kube-system exec -it $(kubectl get pod -n kube-system -l app=antrea-agent -o name) -- ovs-ofctl dump-flows
```

Ingress Controllers & Contour Example

[Ingress purpose](#) – expose HTTP/HTTPS services on a single external IP, routing based on hostnames/paths.

Deploy Contour in kind

```
git clone https://github.com/projectcontour/contour.git
kubectl apply -f contour/examples/contour
```

Add a host entry for local testing:

```
echo "127.0.0.1 my-service.local" | sudo tee -a /etc/hosts
```

Ingress resource for my-service

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example-ingress
spec:
  rules:
  - host: my-service.local
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: my-service
            port:
              number: 8080
```

Figure 4.6 – Ingress routing flow (illustrative; not in transcript)

Note – The upcoming [Gateway API](#) will eventually replace the Ingress API, offering richer routing primitives.

Pod Storage Fundamentals

Definition – *PersistentVolume (PV)*: a cluster-wide storage resource provisioned by an administrator.

Definition – *PersistentVolumeClaim (PVC)*: a request for storage by a pod.

Definition – *StorageClass*: a profile that determines how a PV is dynamically provisioned (e.g., standard, local-path).

Example PVC & Pod (dual-volume)

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: dynamic1
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100Ki
---
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: busybox
      image: busybox
      volumeMounts:
        - name: shared
          mountPath: /shared
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
      volumeMounts:
        - name: dynamic1
          mountPath: /var/www
        - name: shared
          mountPath: /shared
  volumes:
    - name: dynamic1
      persistentVolumeClaim:
        claimName: dynamic1
    - name: shared
      emptyDir: {}
```

Result: busybox and nginx share /shared (ephemeral emptyDir), while nginx writes persistent data to /var/www backed by the dynamically provisioned PVC.

Verifying PVC provisioning

```
kubectl get sc
# Example output:
# NAME      PROVISIONER      RECLAIMPOLICY   AGE
# standard (default)  rancher.io/local-path  Delete   9d
```

📌 Summary of Core Concepts

- **Prometheus** provides gauges, histograms, and counters to observe cluster health and pinpoint outages.
- **kube-proxy** translates Service IPs to Pod IPs using iptables/IPVS; it is distinct from the CNI data plane.
- **CNI plugins** (Calico, Antrea, etc.) attach pods to a network, each employing different routing (BGP vs. OVS).

- **NetworkPolicies** enforce pod-level firewalls; Calico uses iptables, Antrea uses OVS flows.
- **Ingress controllers** (e.g., Contour) map external HTTP traffic to internal Services; the future **Gateway API** will extend this model.
- **Persistent storage** in Kubernetes is abstracted via PV, PVC, and StorageClass, allowing pods to retain data across restarts.

These notes are modular; each ## section can be combined with other study guides as needed.##  Dynamic Storage Provisioning

Dynamic provisioning automatically creates a PersistentVolume (PV) when a PersistentVolumeClaim (PVC) is submitted, without manual PV creation.

- The controller watches for PVCs, calls the provisioner (e.g., **local-path-provisioner** bundled with kind), and creates a PV that is immediately bound to the PVC.
- After binding, the scheduler confirms the PVC is deployable, moving the Pod from **Pending** → **ContainerCreating** → **Running**.

```
$ kubectl get pv
NAME          CAPACITY   ACCESS MODES  RECLAIM POLICY  STATUS   CLAIM
pvc-74879bc4-e2da-4436-9f2b-5568bae4351a  100k       RWO           Delete        Bound    default/dynam
```

The standard **StorageClass** (see image below) tells Kubernetes how to provision the volume.

```
$ kubectl get sc -o yaml
apiVersion: v1
items:
- apiVersion: storage.k8s.io/v1
  kind: StorageClass
  metadata:
    annotations:
      kubectl.kubernetes.io/last-applied-configuration: |
        {"apiVersion":"storage.k8s.io/v1",
         "kind":"StorageClass","metadata":{{
           "annotations":{{
             "storageclass.kubernetes.io/is-default-class": "true"
             , "name": "standard"
           }},
           "provisioner":"rancher.io/local-path",
           "reclaimPolicy": "Delete",
           "volumeBindingMode": "WaitForFirstConsumer"
         }}}
        storageclass.kubernetes.io/is-default-class: "true"
      name: standard
      provisioner: rancher.io/local-path
  kind: List
```

The **is-default-class** makes this the go-to volume for Pods wanting storage without needing to explicitly request a storage class.


You can have many different storage classes in a cluster.

The YAML defines a standard StorageClass with the `rancher.io/local-path` provisioner, a `Delete` reclaim policy, and `WaitForFirstConsumer` binding mode. Declaring a class as `default` lets PVCs without an explicit class automatically receive it.

Container Storage Interface (CSI)

CSI provides a vendor-agnostic plug-in model for storage, replacing the older *in-tree* drivers baked into the kubelet.

- Vendors ship CSI drivers as DaemonSets (node side) and controllers (control-plane side).
- CSI decouples storage logic from the Kubernetes core, allowing independent updates of storage software.



The diagram visualizes the interaction between the kubelet, CSI controller, node plugin, and the container runtime during pod startup.

CSI Architecture & Services

| Service Category | Purpose | Typical RPC Methods |
|-------------------|--|---|
| Identity | Self-identifies the driver (metadata, version) so the control plane knows it is available. | GetPluginInfo, GetPluginCapabilities |
| Controller | Handles volume lifecycle: create, delete, attach, detach. | CreateVolume, DeleteVolume, ControllerPublishVolume |
| Node | Executes node-local actions such as mounting and unmounting volumes. | NodeStageVolume, NodePublishVolume, NodeUnstageVolume |

These services communicate over [gRPC](#) sockets (e.g., `/plugin/csi.sock`) that the kubelet contacts.

```
// Simplified CSI service definitions (highlighted for clarity)
service Identity {
    rpc GetPluginInfo (GetPluginInfoRequest) returns (GetPluginInfoResponse) {}
}
service Controller {
    rpc CreateVolume (CreateVolumeRequest) returns (CreateVolumeResponse) {}
}
service Node {
    rpc NodePublishVolume (NodePublishVolumeRequest) returns (NodePublishVolumeResponse) {}
}
```

```

service Identity {
    rpc GetPluginInfo(GetPluginInfoRequest) <--> The Identity service tells Kubernetes
    rpc GetPluginCapabilities(GetPluginCapabilitiesRequest) what type of volumes can be created by
    rpc Probe (ProbeRequest) the controllers running in a cluster.

}

service Controller {
    rpc CreateVolume (CreateVolumeRequest) <--> The Create and Delete methods are
    rpc DeleteVolume (DeleteVolumeRequest) called before a node can mount a
    rpc ControllerPublishVolume (ControllerPublishVolumeRequest) volume into a Pod, implementing
} dynamic storage.

service Node {
    rpc NodeStageVolume (NodeStageVolumeRequest) <--> The Node service is the part of CSI
    rpc NodeUnstageVolume (NodeUnstageVolumeRequest) that runs on a kubelet, mounting
    rpc NodePublishVolume (NodePublishVolumeRequest) the volume created previously into
    rpc NodeUnpublishVolume (NodeUnpublishVolumeRequest) a specific Pod on demand.
    rpc NodeGetInfo (NodeGetInfoRequest)
}
...
}

```

The snippet shows the three CSI services and their key RPCs, illustrating how a storage vendor implements the specification.

CSI Driver Workflow

1. **Register driver** – The driver advertises its name and capabilities via the **Identity** service.
2. **CreateVolume** – The controller (external provisioner) receives a PVC event, calls CreateVolume on the driver, and receives a volume ID.
3. **PublishVolume** – The scheduler binds the PVC to a node; the node plugin receives a NodePublishVolume request and performs the bind mount.
4. **Mount** – The kubelet mounts the volume into the pod's filesystem namespace.

Bind mount is the Linux operation that mirrors a directory into another location, enabling the pod to access the provisioned storage.

Storage Types Comparison

| Storage Type | Lifespan | Durability | Typical Use-Case |
|----------------------|--|--------------------------------|---|
| Local | Persistent on the node | High (disk stays) | Heavy-weight, data-intensive apps that need node-local fast storage |
| emptyDir | Exists while the pod runs on the node | None (lost on pod termination) | Scratch space, fast temporary storage, shared logs |
| PVC (dynamic) | Persistent across nodes (depending on class) | Configurable (Retain/Delete) | General purpose storage, stateful workloads, databases |

EmptyDir is cheap and fast (can be tmpfs), but unsuitable for data that must survive pod restarts.

PersistentVolume & PVC Lifecycle

1. **Pod request** – A Pod manifest includes a volumeClaimTemplates (StatefulSet) or a direct PVC reference.
2. **Scheduler** – Finds a node meeting the storage topology.
3. **PVC creation** – The control plane watches PVC objects; the appropriate controller creates a PV via the CSI driver.
4. **Binding** – The PV is bound to the PVC; the scheduler now considers the pod runnable.
5. **Kubelet mounts** – Using CSI node services, the volume is mounted into the pod's container.
6. **Pod runs** – Containers read/write to the mounted volume.

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mydata
spec:
  storageClassName: fast
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
```

StorageClasses & Binding Modes

- [StorageClass](#) abstracts the provisioner details (e.g., rancher.io/local-path, pd-ssd).
- [BindingMode](#) determines when the volume is provisioned:
 - Immediate – volume created as soon as PVC appears.
 - WaitForFirstConsumer – provisioning deferred until the pod is scheduled, improving data locality.

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: fast
provisioner: rancher.io/local-path
reclaimPolicy: Delete
volumeBindingMode: WaitForFirstConsumer
parameters:
  type: pd-ssd
```

Admins can define multiple classes (e.g., `bigdata`, `postgres`, `ai`) to match workload performance and retention requirements.

HostPath & CNI Integration

- [hostPath](#) mounts a directory from the node into a pod. Useful for system-level tasks (e.g., CNI plugins installing binaries).
- Many CNI and CSI drivers rely on hostPath for privileged operations such as writing binaries to `/opt/cni/bin` or exposing host sockets.

```
volumes:
- name: cni-bin
  hostPath:
    path: /opt/cni/bin
    type: DirectoryOrCreate
```

Caution: Using hostPath is often considered an anti-pattern for application workloads because it ties pods to specific nodes.

CSI on Non-Linux OS

- CSI is OS-agnostic; on Windows, the [csi-proxy](#) project mediates between the kubelet and Windows-specific storage APIs (PowerShell, etc.).
- As Windows support matures, CSI drivers can run as DaemonSets on Windows nodes, similar to Linux deployments.

Summary

- [Dynamic provisioning](#) + [CSI](#) = fully automated, vendor-agnostic storage that scales with workloads.

- **StorageClasses** let administrators expose multiple storage profiles (performance, retention, locality) without changing application YAML.
- **emptyDir** provides fast, transient storage; **local** offers node-bound durability; **PVC** (via CSI) delivers persistent, portable volumes.
- The **CSI** model separates **identity**, **controller**, and **node** services, simplifying driver development and enabling hot-swap of storage back-ends.
- **hostPath** remains a powerful tool for system-level integrations (CNI/CSI bootstrapping) but should be avoided for regular app data.
- CSI works across Linux and Windows, with platform-specific helpers (e.g., csi-proxy) handling OS differences.

These notes can be combined with other sections to form a complete Kubernetes storage study guide.##  DNS & Pod Networking

DNS resolution in a cluster links hostnames to IPs, which are then routed by the network proxy or CNI.

- **Service IP** → network proxy forwards to a pod endpoint.
- **Pod IP** → CNI makes the IP directly routable.
- **External IP** → outbound traffic is NAT-ed via iptables so responses return to the originating pod.

/etc/resolv.conf inside a pod (example):

```
search default.svc.cluster.local svc.cluster.local cluster.local
nameserver 10.96.0.10
options ndots:5
```

- The **search** list appends domains until a query succeeds.
- **Nameserver 10.96.0.10** is the cluster-internal **kube-dns/CoreDNS** service.

Example DNS query (from an NGINX pod):

```
# wget web-ss-0.nginx.default.svc.cluster.local
Connecting to web-ss-0.nginx.default.svc.cluster.local (10.244.0.13:80)
```

CoreDNS – Cluster-First Resolver

CoreDNS processes DNS queries through a chain of **plugins**, each on its own line in the ConfigMap.

```
kubectl get cm coredns -n kube-system -o yaml
```

```

apiVersion: v1
data:
  Corefile: |
    .:53 {
      errors
      health {
        lameduck 5s
      }
      ready
      kubernetes
        cluster.local in-addr.arpa ip6.arpa {
          pods insecure
          fallthrough in-addr.arpa ip6.arpa
          ttl 30
        }
      prometheus :9153
      forward . /etc/resolv.conf
      cache 30
      loop
      reload
      loadbalance
      log {
        class all
      }
    }
kind: ConfigMap

```

The image shows a CoreDNS ConfigMap with plugins for health checks, Kubernetes integration, and logging.

- **Kubernetes plugin** resolves cluster-local names.
- If it fails, the **forward** plugin uses the host's resolv.conf to reach the internet.

Pod DNS policy

| dnsPolicy | Effect |
|------------------------|---|
| ClusterFirst (default) | Pods use CoreDNS as primary resolver. |
| Default | Pods inherit the node's /etc/resolv.conf. |

⚙️ Tuning CoreDNS Cache

The **cache** plugin stores results (default 30s). Long cache times can delay DNS updates after scaling events.

Fix:

```
kubectl edit cm coredns -n kube-system # set cache TTL to 5s
```

| Parameter | Default | Recommended (small clusters) |
|-----------|---------|------------------------------|
| cache TTL | 30s | 5s (faster refresh) |

🏗️ Control Plane Overview (Kind)

Kind runs a full control plane inside Docker containers.

```
kind create cluster
kubectl -n kube-system get po -o custom-columns=:metadata.name"
```

```

$ docker exec -it \
$(docker ps | grep kind | awk '{print $1}') \
/bin/bash
root@kind-control-plane:/# ps aux | grep \
"/usr/bin/kubelet"
root    722 11.7  3.5 1272784 71896 ?      Ssl  23:34   ← ps (process status) for
➥ 1:10 /usr/bin/kubelet
➥ --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf
➥ --kubeconfig=/etc/kubernetes/kubelet.conf
➥ --config=/var/lib/kubelet/config.yaml --container-runtime=remote
➥ --container-runtime-endpoint=/run/containerd/containerd.sock
➥ --fail-swap-on=false --node-ip=172.17.0.2
➥ --fail-swap-on=false           ← The running kubelet process

```

The screenshot lists core control-plane pods (`coredns`, `etcd`, `kube-apiserver`, etc.).

Key components:

- `etcd` – strongly consistent key/value store.
 - `kube-apiserver` – RESTful front-end.
 - `kube-controller-manager` – runs controllers (node, replication, endpoints).
 - `kube-scheduler` – assigns pods to nodes.
 - `kube-proxy` – service networking.
-

API Server & Resources

The API server exposes built-in and custom resources.

```
kubectl api-resources -o name | head -n 20
```

```

$ kubectl api-resources -o name | head -n 20   ← Shows the available APIs,
bindings                                         just looking at the first 20
componentstatuses
configmaps
endpoints
events
limitranges
namespaces
nodes
persistentvolumeclaims
persistentvolumes
pods
podtemplates
replicationcontrollers
resourcequotas
secrets
serviceaccounts
services
mutatingwebhookconfigurations.admissionregistration.k8s.io
validatingwebhookconfigurations.admissionregistration.k8s.io
customresourcedefinitions.apiextensions.k8s.io

```

List of core and custom API groups (e.g., `bindings`, `configmaps`, `admissionregistration.k8s.io`).

API versioning

| Level | Stability | Typical use |
|--------------------|---|---|
| <code>Alpha</code> | Experimental, <code>not</code> for production | Early feature testing |
| <code>Beta</code> | Production-ready, guaranteed support | Widely used (e.g., <code>DaemonSet</code>) |

| | | |
|----|--------------|---------------------|
| GA | Fully stable | Default in clusters |
|----|--------------|---------------------|

Custom Resource Definitions (CRDs)

- Enable developers to add [own API objects](#) without modifying the core API server.
- Applied with kubectl apply -f <CRD-yaml>.
- Used by operators, Istio, Envoy, etc., to extend cluster functionality.

Scheduler Plugin Framework

The scheduler reads a [Plugins](#) struct; each phase runs a specific set of plugins.

```
type Plugins struct {
    QueueSort    *PluginSet
    PreFilter   *PluginSet
    Filter      *PluginSet
    PostFilter  *PluginSet
    PreScore    *PluginSet
    Score       *PluginSet
    Reserve     *PluginSet
    Permit      *PluginSet
    PreBind    *PluginSet
    Bind        *PluginSet
    PostBind   *PluginSet
}
```

```
PreScore: &schedulerapi.PluginSet{           ←
    Enabled: []schedulerapi.Plugin{
        {Name: interpodaffinity.Name},
        {Name: podtopologyspread.Name},
        {Name: tainttoleration.Name},
    },
},
```

All the defined plugins have already run during the filtering process.

Snippet shows PreScore plugin list; comment notes that preceding plugins have already run.

Key phases (simplified)

1. **QueueSort** – orders pending pods.
2. **PreFilter** – prepares node data (e.g., resource fit).
3. **Filter** – discards unschedulable nodes (taints, affinity, volume limits).
4. **Score** – ranks remaining nodes (balanced resources, image locality).
5. **Bind** – writes the binding to the API server.

[Scoring example](#) (CPU & memory weighting):

```
[ \text{Score} = \frac{\text{cpu\_fraction}}{\text{weightSum}} + \frac{\text{memory\_fraction}}{\text{weightSum}} ]
```

Controller Manager & Storage

- **Node controller** watches node health and updates the Node API object.
- **Replication controller** (now [ReplicaSet](#)) maintains desired pod counts.
- **Endpoint controller** populates Endpoints objects for services.

[Headless Service](#) (no ClusterIP) returns pod IPs directly; useful for StatefulSets.

Service Accounts & Tokens

- Every namespace gets a [default ServiceAccount](#).
 - Pods inherit it unless serviceAccountName is specified.
 - Token is auto-mounted; disable with automountServiceAccountToken: false.
-

Cloud Controller Manager (CCM)

CCM abstracts cloud-specific actions (nodes, routes, load balancers).

```
type Interface interface {
    Initialize(clientBuilder ControllerClientBuilder, stop <-chan struct{})
    LoadBalancer() (LoadBalancer, bool)
    Instances() (Instances, bool)
    // ... other cloud-specific methods ...
}
```

- Implements three main controllers (node, route, service).
 - Runs as a single binary but may spawn additional control loops.
-

etcd – The Consistent Data Store

etcd provides [strong consistency](#) via the [Raft](#) consensus algorithm.

Raft basics (simplified)

1. One [leader](#) and multiple [followers](#) (odd number of nodes).
2. Client writes → leader → replicates to majority.
3. Once a majority acknowledges, the write is committed.
4. If the leader fails, followers elect a new leader.

Write-ahead log (WAL)

- Guarantees durability: a write returns only after [fsync](#) to disk.
- Performance tip: [fsync≤250ms](#) on SSDs; >1s indicates trouble.

Performance histogram example

| le (seconds) | Count |
|--------------|-------|
| 0.001 | 1239 |
| 0.002 | 2365 |
| ... | ... |
| 1.024 | 2588 |
| +Inf | 2588 |

Most writes complete well under [0.008s](#) in a healthy cluster.

Size limits

| Setting | Default |
|---------------------------|---|
| Max request size | 1.5 MiB (--max-request-bytes) |
| DB size recommendation | ≤8 GiB (typical <2 GiB) |
| Memory for large clusters | ≈64 GiB (to hold watches) |

Encryption at rest

- API server flag --encryption-provider-config points to a YAML defining providers (AES-GCM, AES-CBC, SecretBox).
- First provider encrypts; all providers are tried for decryption.

etcd Operations

```
# Health check (embedded perf test)
etcdctl --endpoints=https://127.0.0.1:2379 \
    --cacert=/etc/kubernetes/pki/etcd/ca.crt \
    --cert=/etc/kubernetes/pki/etcd/server.crt \
    --key=/etc/kubernetes/pki/etcd/server.key \
check perf
```

- **Throughput** ≈ 150 writes/s, **slowest** ≈ 0.2s on a small kind cluster.

Running etcd in a pod for debugging

```
apiVersion: v1
kind: Pod
metadata:
  name: etcdclient
  namespace: kube-system
spec:
  hostNetwork: true
  containers:
  - name: etcdclient
    image: ubuntu
    command: ["sleep", "6000"]
    env:
      - name: ETCDCTL_API
        value: "3"
      - name: ETCDCTL_ENDPOINTS
        value: "https://127.0.0.1:2379"
    volumeMounts:
      - name: etcd-certs
        mountPath: /etc/kubernetes/pki/etcd
  volumes:
  - name: etcd-certs
    hostPath:
      path: /etc/kubernetes/pki/etcd
```

Container & Pod Security

Best practices

1. **Never run containers as root** – define a non-root user (runAsUser).
2. **Use minimal base images** (e.g., Google Distroless).
3. **Scan images** for CVEs (e.g., Trivy, Clair).
4. **Apply linters** (hadolint, shellcheck) in CI pipelines.

SecurityContext example

```

apiVersion: v1
kind: Pod
metadata:
  name: sc-Pod
spec:
  securityContext:
    runAsUser: 3042
    runAsGroup: 4042
    fsGroup: 5042
    fsGroupChangePolicy: "OnRootMismatch"
  volumes:
  - name: sc-vol
    emptyDir: {}
  containers:
  - name: sc-container
    image: my-container
    volumeMounts:
    - name: sc-vol
      mountPath: /data/foo

```

When the Pod starts, NGINX runs with user ID 3042.

The user ID 3042 belongs to group 4042.

If the NGINX process writes any files, they are written with the group ID of 5042.

Changes the ownership of the volume before mounting the volume to the Pod

A mount point

The snippet sets `runAsUser: 3042`, `runAsGroup: 4042`, `fsGroup: 5042`, and mounts an `emptyDir` volume.

```

spec:
  securityContext:
    runAsUser: 3042
    runAsGroup: 4042
    fsGroup: 5042

```

Adding Linux capabilities

```

apiVersion: v1
kind: Pod
metadata:
  name: net-cap
spec:
  containers:
  - name: net-cap
    image: busybox
    securityContext:
      runAsUser: 3042
      runAsGroup: 4042
      fsGroup: 5042
    capabilities:
      add: ["NET_ADMIN", "BPF"]

```

Gives the user `CAP_NET_ADMIN` and `CAT_BPF` capabilities

Pod adds `NET_ADMIN` and `BPF` capabilities without granting full root.

```

securityContext:
  capabilities:
    add: ["NET_ADMIN", "BPF"]

```

Pod Security Policies (PSP) (*deprecated* → *Pod Security Admission*)

```

apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: restricted
spec:
  privileged: false
  runAsUser:
    rule: MustRunAsNonRoot

```

```

seLinux:
  rule: RunAsAny
volumes:
  - 'configMap'
  - 'secret'

```

- Enforced at admission time; blocks non-compliant pods.

ServiceAccount token automount

```

apiVersion: v1
kind: Pod
metadata:
  name: no-sa
spec:
  automountServiceAccountToken: false

```

- Disables the default token for pods that don't need API access.

Summary Tables

CoreDNS Cache Impact

| Action | Expected DNS latency | Observation with 30s cache |
|-----------------------------------|----------------------|---|
| Scale-down → Scale-up StatefulSet | <1s | May experience several-second stalls until cache expires. |
| Reduce TTL to 5s | <1s | Faster DNS refresh, lower stale-record risk. |

etcd Performance Checklist

| Check | Command | Success criteria |
|---------------|---|----------------------------------|
| Disk latency | curl localhost:2381/metrics | grep fsync` |
| Leader health | etcdctl endpoint status --write-out=table | One leader, others follower |
| Encryption | Verify --encryption-provider-config present | Secrets stored encrypted at rest |

Key Takeaways

- **DNS** is tightly coupled with pod networking; resolv.conf points to CoreDNS (10.96.0.10).
- **CoreDNS** uses plugins; tuning cache TTL improves DNS freshness.
- The **control plane** (etcd, apiserver, scheduler, controller-manager) can be inspected with kind.
- **API resources** are versioned (alpha → beta → GA); CRDs extend the API without core changes.
- **Scheduler** operates via a plugin pipeline (QueueSort → Filter → Score → Bind).
- **etcd** provides strict consistency via Raft; monitor fsync latency and leader elections.
- **Security** starts at the container: avoid root, use minimal images, scan for CVEs, apply security contexts, limit capabilities, and disable unnecessary service-account tokens.

These notes can be combined with other sections to build a complete Kubernetes study guide.##  TLS Cipher Suite Management

- **Protocol:** Transport Layer Security ([TLS](#))
- **Key exchange:** Elliptic Curve Diffie-Hellman Ephemeral ([ECDHE](#))
- **Authentication:** Elliptic Curve Digital Signature Algorithm ([ECDSA](#))
- **Encryption:** AES-256 in CBC mode

Note: Monitor your TLS posture to ensure it aligns with organizational standards.

- Update cipher suites on any Kubernetes service via the `--tls-cipher-suites` flag, e.g.:

```
--tls-cipher-suites=TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA,TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA256
```

- Apply the same change to [API server](#), [scheduler](#), [controller-manager](#), and [kubelet](#) if a suite is vulnerable.
- Avoid over-restricting suites; some clients (e.g., AWS ELB) may not support newer elliptic-curve ciphers. Consider [TCP health checks](#) instead of HTTPS when using custom cipher suites.

💻 Immutable OS vs. Patching Nodes

- Immutable OS:** All binaries are read-only; nodes are replaced rather than patched.
- Benefits:
 - Eliminates on-disk modifications → reduces attack surface.
 - Simplifies node replacement & version upgrades.
- Kubernetes components on a [control-plane node](#):

| Component | Present |
|--------------------------------|---------|
| kubelet binary | ✓ |
| kube-proxy binary | ✓ |
| Container runtime (containerd) | ✓ |
| etcd image | ✓ |

- Worker node:** Same as above [except](#) etcd (not supported on Windows).
- For Windows workers, build custom immutable images because the OS is not redistributable.
- See [Tanzu Community Edition](#) for a ready-made immutable-image workflow.

🛡 Isolated Container Runtimes

- [Docker Engine](#) uses namespaces but still shares the host kernel → vulnerable if the kernel is compromised.
- Projects providing a [virtual kernel sandbox](#):

| Project | Status | Used by |
|-----------------------------|----------|--------------|
| gVisor | Mature | Google Cloud |
| Firecracker | Mature | AWS |
| Kata | Mature | Various |
| IBM Nabla | Emerging | - |

- These runtimes run containers on a [micro-VM](#) or [virtualized kernel](#), offering stronger isolation.

⚡ Resource Attacks & Limits

Why limits matter

- Pods share finite CPU, memory, and disk on a node.
- Unrestricted containers can cause a [Denial-of-Service \(DoS\)](#) on the node.

Resource limit YAML example

```
apiVersion: v1
kind: Pod
metadata:
  name: core-kube-limited
spec:
  containers:
    - name: app
      image: my-image
      resources:
        requests:
          memory: "42Mi"
          cpu: "42m"
          ephemeral-storage: "21Gi"
        limits:
          memory: "128Mi"
          cpu: "84m"
          ephemeral-storage: "42Gi"
```

- **Requests** → minimum resources the scheduler reserves.
- **Limits** → maximum resources; exceeding triggers a restart, then termination if repeated.

CPU Units

Kubernetes treats **1 CPU** as one hyper-thread (bare metal) or one vCPU (cloud).

| Notation | Meaning |
|----------|-------------------------------|
| 1 | 1 full CPU core |
| 0.25 | $\frac{1}{4}$ core |
| 250m | 250 millicores (same as 0.25) |

```
resources:
  requests:
    cpu: "42"
```

Sets 42 CPUs (it's a big server!)

```
resources:
  requests:
    cpu: "0.42"
```

0.42 of a CPU that is measured as a unit of 1

```
resources:
  requests:
    cpu: "420m"
```

This is the same as 0.42 in the previous code block.

The image illustrates three CPU request styles—42 CPUs (big server), 0.42 CPUs, and 420m—showing how the same capacity can be expressed differently.

Memory Units

Memory can be specified using decimal or binary suffixes:

| Suffix | Decimal | Binary |
|--------|--------------|----------------|
| K / Ki | 10^3 bytes | 2^{10} bytes |
| M / Mi | 10^6 bytes | 2^{20} bytes |
| G / Gi | 10^9 bytes | 2^{30} bytes |
| ... | ... | ... |

Examples (all ~129 MB):

```
resources:
  requests:
    memory: "128974848"      # plain bytes
    memory: "129e6"           # scientific notation
    memory: "129M"            # megabits ≈ 1.613 × 107 B
    memory: "123Mi"           # mebibytes
```

Ephemeral Storage

- Applies to [emptyDir](#) (except tmpfs), node-level log directories, and writable container layers.
- Exceeding limits → [kubelet evicts](#) the Pod.

Extended resources (e.g., GPUs) have separate limits defined in the Kubernetes docs.

Host Networks vs. Pod Networks

- [Pod network \(default\)](#) → isolates Pods from host ports; recommended for security.
- [Host network](#) (hostNetwork: true) → Pod shares the node's network namespace → larger [blast radius](#).

```
apiVersion: v1
kind: Pod
metadata:
  name: host-Pod
spec:
  hostNetwork: true
```

Tip: Use host networking only for system-level Pods (e.g., CNI agents), never expose them directly to the Internet.

Full-Featured Pod Example

```
apiVersion: v1
kind: Pod
metadata:
  name: example-Pod
spec:
  automountServiceAccountToken: false
  securityContext:
    runAsUser: 3042
    runAsGroup: 4042
    fsGroup: 5042
    capabilities:
      add: ["NET_ADMIN"]
  hostNetwork: true
  volumes:
  - name: sc-vol
```

```

emptyDir: {}
containers:
- name: sc-container
  image: my-container
  resources:
    requests:
      memory: "42Mi"
      cpu: "42m"
      ephemeral-storage: "1Gi"
    limits:
      memory: "128Mi"
      cpu: "84m"
      ephemeral-storage: "2Gi"
  volumeMounts:
  - name: sc-vol
    mountPath: /data/foo
serviceAccountName: network-sa

```

- Disables default service-account token.
- Sets explicit UID/GID and adds NET_ADMIN.
- Runs on the host network with resource limits.

API Server Security & RBAC

Admission Webhooks

- [ImagePolicyWebhook](#): validates container images before admission.
- Pods failing admission stay [Pending](#).

Role-Based Access Control (RBAC)

- Enabled via --authorization-mode=RBAC.
- Core RBAC objects:

| Object | Scope |
|------------------------------------|--|
| Role | Namespace-limited |
| ClusterRole | Cluster-wide |
| RoleBinding | Grants a Role to users/groups/service-accounts |
| ClusterRoleBinding | Grants a ClusterRole globally |

Verbs (allowed actions)

- API resource verbs: get, list, create, update, patch, watch, proxy, redirect, delete, deletecollection
- Non-resource HTTP verbs: get, post, put, delete

Example Role & Binding (list & patch Pods)

```

# Role
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: Pod-labeler
  namespace: rbac-example
rules:
- apiGroups: [""]
  resources: ["Pods"]
  verbs: ["list", "patch"]
---
# ServiceAccount
apiVersion: v1

```

```

kind: ServiceAccount
metadata:
  name: Pod-labeler
  namespace: rbac-example
---
# RoleBinding
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: Pod-labeler
  namespace: rbac-example
subjects:
- kind: ServiceAccount
  name: Pod-labeler
  roleRef:
    kind: Role
    name: Pod-labeler
    apiGroup: rbac.authorization.k8s.io

```

Subresources (e.g., Pod logs)

```

apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: Pod-and-Pod-logs-reader
  namespace: rbac-example
rules:
- apiGroups: [""]
  resources: ["Pods", "Pods/log"]
  verbs: ["get", "list"]

```

Subjects can include [Users](#), [Groups](#), [ServiceAccounts](#).

Debugging RBAC

- [Audit logs](#) capture *who, what, where, how* for every API request.
- Enable via --audit-policy-file on the API server.
- Helpful tools:

| Tool | Purpose |
|----------------------------|--------------------------------|
| rbac-lookup (ReactiveOps) | List roles for a subject |
| kubectl-rbac (OctarineSec) | Show permissions for a subject |
| audit2rbac (Jordan Liggit) | Generate RBAC from audit logs |

Authn, Authz, and Secrets

- Do **not** use default admin certificates.
- Prefer [IAM service accounts](#) for cloud-native identity.
- Disable basic username/password auth; rely on TLS client certificates.

IAM Service Accounts (cloud)

- Every node gets an IAM role; Pods inherit it unless overridden.
- [Least-privilege](#): create dedicated roles for each workload.
- Tools for per-Pod IAM isolation: kube2iam, kiam.

Private API Servers

- Place the API server behind a [private network](#) or bastion host.
 - Whitelist load-balancer IPs; avoid exposing the API server to the Internet.
-

Network Security

NetworkPolicies (requires a CNI that supports them, e.g., Calico)

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-ingress
  namespace: web
spec:
  podSelector: {}
  policyTypes: ["Ingress"]
```

- Denies all inbound traffic to pods in web.

Allow traffic from a specific namespace

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: access-web
  namespace: web
spec:
  podSelector:
    matchLabels:
      app: nginx
  policyTypes: ["Ingress"]
  ingress:
  - from:
    - namespaceSelector:
        matchLabels:
          name: test-bed
```

- [Label namespaces](#) (kubectl label namespaces test-bed name=test-bed) to make the selector work.

Default deny-all (both ingress & egress)

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-all
  namespace: test-bed
spec:
  podSelector: {}
  policyTypes: ["Ingress", "Egress"]
```

- After applying, pods lose DNS access; add an egress rule for DNS if needed:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: dns-egress
  namespace: test-bed
spec:
  podSelector: {}
  policyTypes: ["Egress"]
```

```
egress:
- to:
  - namespaceSelector:
    matchLabels:
      name: kube-system
ports:
- protocol: UDP
  port: 53
```

Service Mesh (brief)

- Provides [mutual TLS](#) and advanced traffic policies.
- Common meshes: [Istio](#), [Linkerd](#), [Consul](#).
- Not required for day-one clusters; consider only if you need its capabilities.

Open Policy Agent (OPA)

- Declarative policy engine usable as an [admission controller](#) or sidecar.
- Two main components:

| Component | Role |
|--|--|
| OPA Gatekeeper | CRD-based policies, audit capabilities |
| OPA admission controller | Direct webhook enforcement |

Installing OPA Gatekeeper (kind example)

```
kind delete cluster
kind create cluster
kubectl apply -f https://raw.githubusercontent.com/open-policy-agent/gatekeeper/v3.7.0/deploy/gatekeeper.yaml
```

```
kubectl -n gatekeeper-system get po
# Expected output: gatekeeper-audit-..., gatekeeper-controller-manager-....
```

Sample ConstraintTemplate (enforce specific container registry)

```
apiVersion: templates.gatekeeper.sh/v1beta1
kind: ConstraintTemplate
metadata:
  name: enforcespecificcontainerregistry
spec:
  crd:
    spec:
      names:
        kind: EnforceSpecificContainerRegistry
  targets:
  - target: admission.k8s.gatekeeper.sh
    rego: |
      package enforcespecificcontainerregistry
      violation[{"msg": msg}] {
        container := input.review.object.spec.containers[_]
        satisfied := [good |
          repo = input.parameters.repos[_];
          good = startswith(container.image, repo)
        ]
        not any(satisfied)
        msg := sprintf("container '%v' has an invalid image repo '%v', allowed repos are %v",
        
```

```
[container.name, container.image, input.parameters.repos])  
}
```

```
apiVersion: constraints.gatekeeper.sh/v1beta1  
kind: EnforceSpecificContainerRegistry  
metadata:  
  name: enforce-specific-container-registry-test-ns  
spec:  
  match:  
    kinds: [{apiGroups: [""], kinds: ["Pod"]}]  
    namespaces: ["test-ns"]  
  parameters:  
    repos: ["myrepo.com"]
```

- Attempting to create a pod with an image outside myrepo.com will be denied.

Multi-Tenancy

| Trust Level | Typical Use-Case |
|------------------|--|
| High trust | Different departments of the same organization sharing a cluster. |
| Medium/Low trust | External customers in separate namespaces. |
| Zero trust | Data subject to legal constraints; require strict isolation (often separate clusters). |

- Kubernetes [namespace](#) is the primary isolation boundary, but it does [not](#) prevent cross-namespace API access.
- For strong isolation, consider [separate clusters](#), [RBAC](#), [NetworkPolicies](#), and possibly a [service mesh](#).

Kubernetes Tips (quick checklist)

- Private API server endpoint (no internet exposure).
- Enable [RBAC](#).
- Apply [NetworkPolicies](#).
- Disable username/password auth.
- Use dedicated service accounts; disable default token mounting.
- Avoid hostNetwork unless required.
- Define [resource quotas](#) and per-Pod limits.

Application Management – Guestbook Example

Installing the classic Guestbook (single YAML)

```
kubectl create -f https://github.com/kubernetes/examples/blob/master/guestbook/all-in-one/guestbook-all-in-one.yaml
```

- [Caution](#): Curling YAML directly from the internet can be risky; prefer trusted repos.

Challenges with monolithic YAML

- Hard to customize, version, or scale.
- Repeating large blocks leads to maintenance overhead.

Modularizing with [Carvel tools](#) ([ytt](#), [kapp](#), [kapp-controller](#))

Install Carvel toolkit

```
# macOS
brew tap vmware-tanzu/carvel && brew install ytt kblf kapp imgpkg kwt vendir
# Linux
curl -L https://carvel.dev/install.sh | bash
```

Split resources into separate files (example directory layout)

```
v1/
fe-dep.yaml
fe-svc.yaml
redis-master-dep.yaml
redis-master-svc.yaml
redis-slave-dep.yaml
redis-slave-svc.yaml
```

Patch CPU requests with `ytt` overlays

Image 2 shows a `ytt` overlay that matches the frontend-dep pod and updates its CPU request to 200m.

```
#@ load("@ytt:overlay", "overlay")
#@overlay/match by=overlay.subset(
  {"metadata": {"name": "frontend-dep"}}) ← Our ytt overlay identifies the
---                                         name of a YAML snippet we
spec:                                         want to match.
  template:
    spec:
      containers:
        #@overlay/match by="name"   ← Once inside of our containers,
        - name: php-redis          it substitutes the container
          #@overlay/match missing_ok=False
          resources:
            requests:
              cpu: 200m             ← The original CPU value of 100m
                                         is now doubled to 200m.
```

The overlay matches `metadata.name: frontend-dep` and substitutes the container's CPU request.

```
# v2/ytt-cpu-overlay.yml
#@ load("@ytt:overlay", "overlay")
#@overlay/match by=overlay.subset({"metadata": {"name": "frontend-dep"}})
---
spec:
  template:
    spec:
      containers:
        #@overlay/match by="name"
        - name: php-redis
          resources:
            requests:
              cpu: 200m # new request
```

- Apply overlays and deploy with `kapp`:

```
kapp deploy -a guestbook \
-f <(ytt -f ./v1/ -f ./v2/ytt-cpu-overlay.yml -f ./v2/ytt-cpu-overlay-db.yml)
```

Image 3 demonstrates the `kapp deploy` command used to install the Guestbook app.

```

$ kapp deploy -a kc -f https://github.com/vmware-tanzu/
  carvel-kapp-controller/releases/latest/
  download/release.yml
$ kapp deploy -a default-ns-rbac -f
  https://raw.githubusercontent.com/vmware-tanzu/
  carvel-kapp-controller/
  develop/examples/rbac/default-ns.yml

```

The diagram shows two lines of command-line text. The first line installs the kapp-controller using the kapp tool, as indicated by a callout pointing to the 'kapp' command. The second line installs a simple RBAC definition, as indicated by a callout pointing to the URL 'https://raw.githubusercontent.com/vmware-tanzu/carvel-kapp-controller/develop/examples/rbac/default-ns.yml'.

The image shows the two-step kapp deployment workflow.

- [Inspect](#) the deployed app:

```
kapp list
kapp inspect --app=guestbook
```

- [Delete](#) atomically:

```
kapp delete --app=guestbook
```

Managing the Guestbook with [kapp-controller](#) (Operator-style)

```

apiVersion: kappctrl.k14s.io/v1alpha1
kind: App
metadata:
  name: guestbook
  namespace: default
spec:
  serviceAccountName: default-ns-sa
  fetch:
    - git:
        url: https://github.com/jayunit100/k8sprototypes
        ref: origin/master
        subPath: carvel-guestbook/v2/output/
  template:
    - ytt: {}
  deploy:
    - kapp: {}

```

[Image 4](#) shows the YAML snippet for a kapp App CR.

```

apiVersion: kappctrl.k14s.io/v1alpha1
kind: App
metadata:
  name: guestbook
  namespace: default
spec:
  serviceAccountName: default-ns-sa
  fetch:
    - git:
        url: https://github.com/jayunit100/k8sprototypes
        ref: origin/master
        # We have a directory, named 'app', in the root of our repo.
        # Files describing the app (i.e. pod, service) are in that directory.
        subPath: carvel-guestbook/v1/
  template:
    - ytt: {}
  deploy:
    - kapp: {}

```

The diagram shows annotations for the kapp App CR YAML snippet. A callout points to the 'serviceAccountName' field with the text 'Uses the service account created previously for this app installation'. Another callout points to the 'fetch' section with the text 'Specifies where our app is defined'. A third callout points to the 'template' section with the text 'Because the code for our app actually lives in carvel-guestbook/v1/, we need to specify this subpath.'

The CR defines where the Git repo lives, which templates to render, and that kapp should perform the deployment.

- Install [kapp-controller](#):

```
kapp deploy -a kc -f https://github.com/vmware-tanzu/carvel-kapp-controller/releases/latest/download/release.  
kapp deploy -a default-ns-rbac -f https://raw.githubusercontent.com/vmware-tanzu/carvel-kapp-controller/devel
```

- Create the Guestbook App CR (`kubectl create -f guestbook-app.yaml`).
- The controller watches the CR, renders YTT templates, and uses Kapp to reconcile the resources—providing [declarative](#), [stateful](#) application management.

Summary Tables

Resource Specification Units

| Resource | Decimal Form | Millicore / Binary Form |
|-------------------|--------------|-------------------------|
| CPU | 0.5 | 500m |
| Memory | 256Mi | 268435456bytes |
| Ephemeral Storage | 10Gi | 10737418240bytes |

Common Security Controls

| Control | Scope | Typical Implementation |
|--------------------------------------|--|---------------------------------------|
| TLS Cipher Suites | API Server, Scheduler, Controller-Manager, Kubelet | --tls-cipher-suites flag |
| Immutable OS | Node OS | Replace VM/image instead of patching |
| Isolated Runtimes | Container runtime | gVisor, Firecracker, Kata |
| Resource Limits | Pods | resources.requests / resources.limits |
| RBAC | API Access | Roles, ClusterRoles, Bindings |
| NetworkPolicies | Pod ⇄ Pod traffic | Calico, Antrea, Cilium |
| OPA/Gatekeeper | Admission control | ConstraintTemplates + Constraints |
| IAM Service Accounts | Cloud API access | Per-Pod roles via kube2iam / kiam |

Further Reading

- [Calico CRD list](#) (example output)

```
kubectl get crd | grep calico  
#@ Output includes bgpconfigurations.crd.projectcalico.org, networkpolicies.crd.projectcalico.org, et
```

 [Image 5](#) displays a fragment of a Calico configuration spec, illustrating the depth of CRD-based configuration.

```

...
spec:
  controllers:
    namespace:
      reconcilerPeriod: 5m0s
    node:
      reconcilerPeriod: 5m0s
      syncLabels: Enabled
    policy:
      reconcilerPeriod: 5m0s
    serviceAccount:
      reconcilerPeriod: 5m0s
    workloadEndpoint:
      reconcilerPeriod: 5m0s
  etcdv3CompactionPeriod: 10m0s
  healthChecks: Enabled
  logSeverityScreen: Info
  prometheusMetricsPort: 9094
status:
  environmentVars:
    DATASTORE_TYPE: kubernetes
    ENABLED_CONTROLLERS: node
runningConfig:
  controllers:
    node:
      hostEndpoint:
        autoCreate: Disabled
        syncLabels: Disabled
    etcdv3CompactionPeriod: 10m0s
  healthChecks: Enabled
  logSeverityScreen: Info

```

We can disable healthChecks if we don't feel it's necessary.

Sets the port on which our Calico kube controller serves up its Prometheus metrics. Here's where we can change the port if needed.

The snippet shows fields such as reconciler period, log severity, and Prometheus metrics port.

- Explore [Tanzu Community Edition](#) for a production-grade, immutable-image Kubernetes distribution.
- Test [NetworkPolicy conformance](#) with Sonobuoy:

```
sonobuoy run --e2e-focus=NetworkPolicy
sonobuoy status
```